

Kaki: Concurrent Update Synthesis for Regular Policies via Petri Games

March 31, 2023


Nicklas S. Johansen, Lasse B. Kær, Andreas L. Madsen,
Kristian Ø. Nielsen, Jiri Srba, and Rasmus G. Tollund

Department of Computer Science
Aalborg University
Kingdom of Denmark



AALBORG UNIVERSITY
DENMARK


Communication Network



- ▶ Switches send packets across a network
- ▶ A flow is a source and destination pair for packets
- ▶ In Software Defined Networking a central controller manages traffic


Network Update

- ▶ Flow from A to C
- ▶ Initial forwarding rules




Network Update

- Flow from A to C
- Initial forwarding rules
- New desired routing



Routing

- ▶ Network
- ▶ Initial Routing
 - ▶ Trace: $s_1 s_2 s_3 s_5 s_6$
- ▶ Final Routing
 - ▶ 4 different traces
- ▶ Intermediate Routing




Policy



- ▶ A policy P is a regular expression over switches V , describing the language $\mathcal{L}(P) \subseteq V^*$
- ▶ A policy P is satisfied for a routing R , if all traces for R is in $\mathcal{L}(P)$


Policy

- ▶ A policy P is a regular expression over switches V , describing the language $\mathcal{L}(P) \subseteq V^*$
- ▶ A policy P is satisfied for a routing R , if all traces for R is in $\mathcal{L}(P)$
- ▶ Policy examples
 - ▶ Reachability




Policy

- ▶ A policy P is a regular expression over switches V , describing the language $\mathcal{L}(P) \subseteq V^*$
- ▶ A policy P is satisfied for a routing R , if all traces for R is in $\mathcal{L}(P)$
- ▶ Policy examples
 - ▶ Reachability
 - ▶ Waypoint




Policy

- ▶ A policy P is a regular expression over switches V , describing the language $\mathcal{L}(P) \subseteq V^*$
- ▶ A policy P is satisfied for a routing R , if all traces for R is in $\mathcal{L}(P)$
- ▶ Policy examples
 - ▶ Reachability
 - ▶ Waypoint
 - ▶ Conditional enforcement ($s_u \Rightarrow s_v$)



Concurrent Update Synthesis Problem

Policy: $V^*\{s_5, s_4\} V^*s_6$




Concurrent Update Synthesis Problem

Policy: $V^*\{s_5, s_4\}V^*s_6$

Concurrent Update Sequence:

- ▶ Batch 1: $\{s_4\}$
- ▶ Batch 2: $\{s_1, s_2, s_3\}$?
 - ▶ Six different permutations
 - ▶ $s_1 \rightarrow s_3 \rightarrow s_2$ **policy violation!**
 - ▶ Violating trace: $s_1 s_2 s_3 s_6 \notin \mathcal{L}(P)$




Concurrent Update Synthesis Problem

Policy: $V^*\{s_5, s_4\}V^*s_6$

Concurrent Update Sequence:

- ▶ Batch 1: $\{s_4\}$
- ▶ Batch 2: $\{s_1, s_2\}$
- ▶ Batch 3: $\{s_3\}$
- ▶ Batch 4: $\{s_5\}$
- ▶ $\omega = \{s_4\}\{s_1, s_2\}\{s_3\}\{s_5\}$







Translation to Petri Game

- ▶ Encode CUSP as Petri game
- ▶ Find winning strategy for controller in Petri game
- ▶ Extract *concurrent update sequence* from strategy

Strategy synthesis using TAPAAL Petri game engine

Petri Game



Translation



CTL query: $AF p3 = 1$

Petri Game



Translation



CTL query: $AF p3 = 1$

Petri Game



Translation



CTL query: $AF p3 = 1$

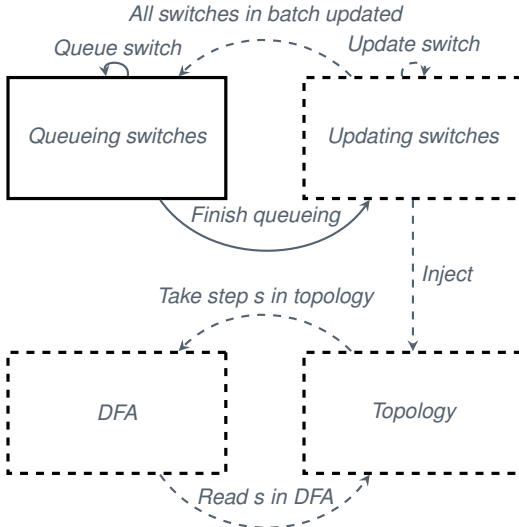
Petri Game

Translation



CTL query: $AF p3 = 1$

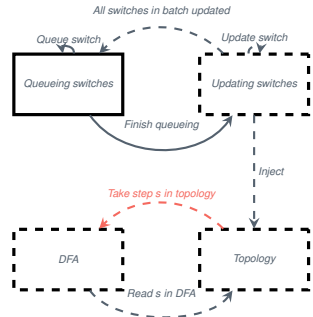

Flowchart for Petri Game



Topology Step

Translation

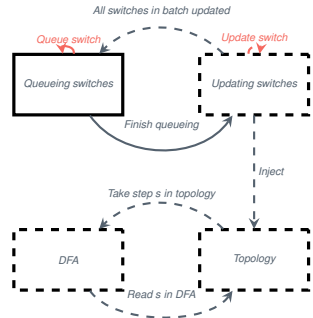

For $s, s' \in V$ where $s' \in R_i(s) \cup R_f(s)$



- A routing from s to s' is encoded by a transition $t_{(s,s')}$
- The policy is encoded as a DFA in the Petri game
- The topology and DFA components are synchronised

Switch Component

Translation



Query Translation




- Solve CUSP by finding a winning strategy for controller

$$AF \left(p_{batches} \leq k \wedge \left(\bigvee_{q \in F} p_q = 1 \right) \right)$$


- Bisection search to find lowest k
- Proved to find optimal solution to CUSP

Topological Decomposition

- Split problem into two independent subproblems
- Find potential NFA states at switches




Waypoint: s_2, s_6




Topological Decomposition

- Split problem into two independent subproblems
- Find potential NFA states at switches




Waypoint: s_2, s_6




Topological Decomposition

- Split problem into two independent subproblems
- Find potential NFA states at switches




Waypoint: s_2, s_6



Switch s_4 is a topological decomposition point!

Subproblems

Topological Decomposition



- Original CUSP is split into 2 smaller CUSPs
- Proven: Combine optimal solution for subproblems into an optimal solution to the original problem

Comparison with FLIP

Results

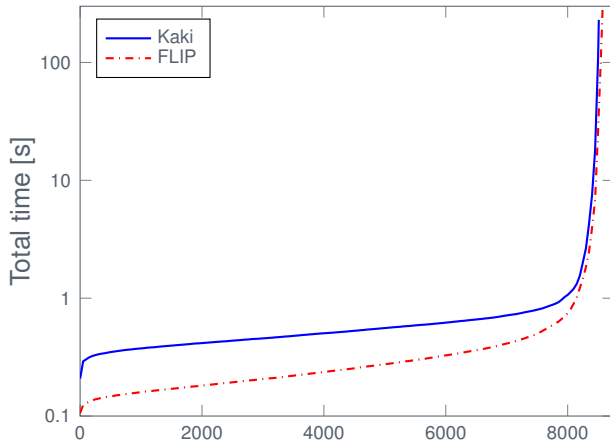


Figure: Kaki benchmarked against FLIP

	All Problems
Total	8759
Only Kaki	133
Only FLIP	196
Suboptimal	787
Tagging	268

Comparison with NetStack

Results




Figure: Comparison with NetStack

Conclusion



- ▶ Kaki capabilities
 - ▶ Find optimal concurrent updates
 - ▶ Supports splittable flows
 - ▶ Supports arbitrary regular policies
- ▶ Performed on par with FLIP

Conclusion



- ▶ Kaki capabilities
 - ▶ Find optimal concurrent updates
 - ▶ Supports splittable flows
 - ▶ Supports arbitrary regular policies
- ▶ Performed on par with FLIP
- ▶ Faster if subproblems are solved parallel

Presentation Concluded

Comparison with SeqPG

Results

- ▶ Kaki limited to 1 switch per batch
- ▶ SeqPG unable to decide 2 problems




Figure: Kaki benchmarked against SeqPG

Splitable Flows

Results

- ▶ Comparison with ourselves
- ▶ Small overhead




Figure: Splitable flows experiments for Kaki

Preprocessing Techniques

Results





Figure: Experiments of the preprocessing techniques of Kaki

Update Equivalence Classes




Always put s_3, s_4, s_8, s_9 in the first batch

Always put s_1, s_2, s_6, s_7 in the last batch

Update Equivalence Classes

Chain reduction



Switches s_3 , s_4 , s_5 can always be updated in the same batch